

The Lab Differentiation Problem

Najmul Hasan




Figure 1: Six frontier labs positioned on openness and compute vertical integration as of May 2026. Each cell pulls together a strategic posture and an infrastructure bet; the empty cell (open weights without a vertically integrated compute stack) is the gap a sufficiently capitalized newcomer could fill.

Six frontier labs train on the same arXiv and recruit from the same labor market. The useful differentiation question, given only public outputs, is what actually separates their training philosophy, research direction, and cultural posture. My read after working through their primary-source output for the last six months is that four axes do nearly all of the distinguishing work. I apply them to Anthropic, OpenAI, Google DeepMind, Amazon AGI, Microsoft AI, and Meta Superintelligence Labs, then specify the methodology for tracking individual researchers' direction from external signals. I have no insider access to any of these labs; every claim is anchored to an original source listed at the end.

Framework

Each axis below is observable from outside the lab and falsifiable against an original source.

- 1. Frontier-model openness.** Closed flagship only, hybrid (closed flagship with open smaller models), or open-weights flagship. Visible on the lab's news page, Hugging Face presence, and license text.
- 2. Compute vertical integration.** Pure cloud tenant, hybrid, or fully integrated stack with custom silicon, owned data centers, and an in-house training framework. Visible from press releases on chip generations and capacity buildouts.
- 3. Safety and governance posture.** Public scaling policy with named capability thresholds, or no public framework. Visible from policy documents on each lab's site and the cadence of per-model safety reports.
- 4. Stated identity.** The mission framing each lab repeats publicly, visible in CEO statements, blog headers, and "about" pages.

Figure 1 positions the six labs on the first two axes; Table 1 gives the full four-axis cut.

Lab	Openness	Compute stack	Safety framework	Stated identity
Anthropic	Closed only (Claude)	Tenant: AWS, GCP, Azure; Trainium partner ([1], [12])	Responsible Scaling Policy ([2])	Safety-first, alignment-led research
OpenAI	Closed flagship + gpt-oss open weights ([4], [5])	Tenant (Azure) shifting to integrated (Stargate, Broadcom, AMD) ([6], [7])	Preparedness Framework v2 ([8])	“AGI that benefits all of humanity”
Google DeepMind	Hybrid (Gemini closed + Gemma 4 open) ([9], [10])	Vertically integrated: TPU Ironwood + Pathways + JAX ([11])	Frontier Safety Framework v3 ([13])	“AI responsibly to benefit humanity,” AI for science
Amazon AGI	Closed (Nova) + Bedrock hosts third-party open models ([14])	Vertically integrated: Trainium3, AWS DCs ([15])	Service cards per model; no scaling policy	Agents that act in digital and physical worlds ([16])
Microsoft AI	Hybrid: MAI closed + Phi open-weights (MIT) + OpenAI partnership ([17], [18])	Vertically integrated: Maia 200, Azure ([19])	Defers to OpenAI Preparedness for frontier; internal RAI standard	“AI companion for everyone” ([20])
Meta SI	Hybrid (Llama 4 open + Muse closed) ([21], [22])	Vertically integrated: MTIA roadmap, Hyperion, nuclear PPAs ([23], [24])	Frontier AI Framework (Meta)	“Personal superintelligence for everyone” ([25])

Table 1: The four-axis cut, each cell anchored to a numbered primary source in the References section.

Per-lab snapshots

Anthropic. Closed weights only, no open releases. Claude Opus 4.7 shipped April 16, 2026 with gains on coding, agents, vision, and multi-step tasks [1]. The published research cadence emphasizes interpretability, alignment, and the Economic Index series. Outside research the recent moves include a \$200M Gates Foundation partnership and the Stainless acquisition on May 18, 2026 [3]. Compute posture is unusual for a frontier lab: Anthropic takes significant capital from Amazon and uses Trainium as a primary training target [12] while also serving on GCP and Azure, which makes them the most cloud-diversified of the six. Every product announcement ships with an explicit safety report.

OpenAI. The most visibly shifting compute posture of the six. Closed flagship with a rapid release cadence (GPT-5.5 announced April 23, 2026, API April 24 [4]) plus the open-weights gpt-oss line under Apache 2.0 [5]. The infrastructure story is bigger than the model story right now. Stargate (\$500B over four years [6]), a 10GW Broadcom co-designed accelerator deployment [7], and a 6GW AMD GPU deal move OpenAI from Azure-tenant to integrated stack on a multi-year timeline. Safety posture is the Preparedness Framework v2, with two capability thresholds and Capabilities plus Safeguards reports [8]. Distinctive research bet: deliberative alignment, training models to explicitly reason over written safety specs [26].

Google DeepMind. Gemini 3.1 Pro is the current flagship [9] with Deep Think variants for hardest reasoning. Gemma 4 ships alongside as open weights [10]; the Gemma ecosystem reports over 400 million downloads. Compute is the most vertically integrated of the six: Ironwood TPUs (the seventh-generation, optimized for inference), Pathways for runtime, JAX for training [11]. The Frontier Safety Framework hit its third iteration in April 2026, introducing Tracked Capability Levels alongside Critical Capability Levels [13]. The scientific-AI line is a continuous program rather than a one-off: AlphaFold 3 with Isomorphic Labs, AlphaEvolve

as a coding agent for research acceleration (introduced May 2025, expanded through 2026), gold-medal IMO 2025 performance from Gemini Deep Think. Hassabis shared the 2024 Nobel in Chemistry (with John Jumper and David Baker) for AlphaFold, and the lab leans into that science identity in its public messaging.

Amazon AGI. Two parallel tracks operate inside Amazon’s AI org. Mainline Nova (Micro, Lite, Pro, Premier) and Nova 2 (Lite GA, Pro preview, December 2025) ship via Bedrock alongside Nova Act for agent-as-a-service [14]. The Amazon AGI SF Lab was formally announced in December 2024 under David Luan (VP Autonomy), built on the Adept team that joined Amazon in June 2024, with an explicit thesis of agents that act in digital and physical worlds [16]. Trainium3 UltraServers reached GA in December 2025 on a 3nm process [15], with Anthropic anchoring as the lead training customer through a cumulative Amazon investment that reached \$8B by late 2024 and has since been expanded [27]. Openness here comes from third-party model choice on Bedrock rather than from open-weights Nova releases.

Microsoft AI. The consumer and Copilot framing is the most explicit of the six. Mustafa Suleyman became EVP and CEO of Microsoft AI in March 2024 with Karen Simonyan as Chief Scientist, importing the Inflection team [17]. The org now ships in-house MAI models (MAI-Voice-1, MAI-Image-2, MAI-Transcribe-1) alongside the Phi small-reasoning family from Microsoft Research [18]. The OpenAI partnership was restructured in October 2025: Microsoft IP rights extended through 2032, AGI declarations now verified by an independent expert panel, an OpenAI commitment to \$250B in additional Azure consumption [28]. The Maia 200 inference chip (3nm) was announced January 26, 2026 [19]. Suleyman frames the work as “an AI companion for everyone” [20], distinct from the AGI framings elsewhere.

Meta Superintelligence Labs. MSL is the youngest of the six and among the most capital-intensive. Alexandr Wang became Meta’s first Chief AI Officer in June 2025 and runs MSL, which houses Llama, products, and FAIR projects [29]. The strategy is hybrid: Llama 4 Scout and Maverick shipped April 5, 2025 as open weights with Behemoth previewed as a teacher model [21], while Muse Spark on April 8, 2026 became the first closed MSL flagship [22], explicitly positioned as “personal superintelligence.” Planned 2026 capex was guided to \$115B–\$135B in Meta’s Q4 2025 8-K [30] and later widened to \$125B–\$145B in the Q1 2026 update, with Hyperion as a Blue Owl JV [24], MTIA custom silicon on a six-month cadence [23], and up to 6.6GW of nuclear PPAs announced January 2026. Zuckerberg’s “personal superintelligence” framing diverges from the AGI framing used by Anthropic, OpenAI, and DeepMind [25].

Reading researcher direction from external signals

Lab-level differentiation is one half of the question. The other half is per-researcher: predicting what a specific person will publish next, from public clues only. A handful of signals carry most of the information.

arXiv preprint patterns. arXiv author listing pages and the cs.LG, cs.CL, cs.AI categories are the leading indicator of where a researcher is heading. arXiv exposes a public API for programmatic queries by author, category, and date [31], which makes longitudinal topic drift directly observable. The current affiliation field in the author block updates faster than any lab’s official announcement page.

Conference acceptance and OpenReview submissions. NeurIPS, ICML, ICLR, and COLM publish accepted-papers lists. OpenReview [32] is uniquely useful because it exposes submissions, reviews, and decisions, so rejected and withdrawn papers are visible. Year-over-year diffs of a lab’s portfolio at a single venue (filtered by affiliation string) show new directions appearing and old ones being abandoned.

Co-authorship graphs. OpenAlex [33] and Semantic Scholar both expose institution-disambiguated

author graphs through open APIs. Louvain communities on the co-authorship graph weighted by joint-paper count over a two-year window cluster cleanly into labs even when affiliation strings are stale. The institution disambiguation pipeline behind OpenAlex is documented in Priem, Piwowar, and Orr (2022) [34].

Researcher public statements. Verified X accounts and personal sites are primary sources when quoted directly. The relevant published methodology is Tang et al. (2008) on ArnetMiner [35] for researcher-profile extraction, and Prabhakaran et al. (2016) on predicting topic trajectories from rhetorical framing [36].

Departures as signal. Major moves between the six labs are reliably timestamped by either the sending lab’s official announcement or the researcher’s first-person post on X. Verifiable public examples from the 2024 to 2026 window: Jan Leike (OpenAI to Anthropic, May 28, 2024 [37]), John Schulman (OpenAI to Anthropic, August 5, 2024 [38]), Ilya Sutskever (OpenAI to Safe Superintelligence Inc., June 19, 2024 [39]), David Luan and the Adept team (to Amazon AGI, June 28, 2024 [40]), Pieter Abbeel and Covariant (to Amazon, August 30, 2024 [41]), Mustafa Suleyman (Inflection to Microsoft, March 19, 2024 [17]), Noam Shazeer and Daniel De Freitas (Character.AI to Google DeepMind, August 2024 [42]), Alexandr Wang (Scale AI to Meta, June 12, 2025 [43]). From watching these moves closely, my observation is that a researcher’s first paper at the new lab usually shows up within a year of the announcement and tends to continue their prior research thread rather than pivot to a new topic. That makes the new arXiv affiliation a usable predictor in practice, not a proven regularity.

GitHub commit activity. Public commits on official lab organizations (anthropics, openai, google-deepmind, meta-llama, amzn, microsoft) precede paper releases by months. Engineering activity reveals direction before the publication does.

Twelve-month forecasts the framework makes

These follow directly from the axes and from your “Kings and Electors” framing on which labs are actually shipping versus being kingmade. Anthropic stays cloud-diversified and refuses to vertically integrate. The Amazon and GCP relationships are deep enough that building their own silicon would mean fighting the relationships funding them. Expect continued differentiation through safety reports and interpretability rather than through compute.

OpenAI is physically different by end of 2026. Stargate, Broadcom, and AMD push the lab from cloud-tenant to integrated stack on a multi-year arc, the same arc DeepMind walked with TPU. Expect more accelerator co-design announcements and reduced dependence on Azure as a platform.

DeepMind compounds on scientific AI. AlphaEvolve, AlphaFold 3, AI Co-Clinician, and the DOE Genesis partnership are not one-offs; science is the cultural moat. Expect another high-visibility scientific result before end of 2026, plus continued Gemma releases that prevent any open-weights challenger from owning that segment.

Amazon AGI continues emphasizing agents-that-act. The SF Lab thesis (originally under David Luan, now under broader Amazon AGI leadership after Luan’s February 2026 departure) is consistent with Nova Act and the lab’s recent blog cadence on browser-use agents and trajectory rollouts. Expect a Nova 3 with explicit agentic capabilities timed near re:Invent 2026.

Microsoft AI sustains the Copilot consumer framing while Microsoft Research carries the technical research load (Phi reasoning models). Expect more MAI-branded product surfaces and continued OpenAI partnership friction surfacing through the periodic restructuring announcements.

Meta SI continues the hybrid Llama-open plus Muse-closed split. The capex envelope (\$125B to

\$145B in 2026 per the Q1 update) is too large to read as a hedge; Muse is the actual frontier bet, while Llama remains the talent magnet and developer-mindshare lever.

The signals are public; the bet I would take from outside is that two of the six labs (DeepMind and Meta SI) materially widen their compute lead through 2026 while Anthropic and OpenAI converge on similar safety frameworks and diverge on infrastructure, with Amazon and Microsoft riding their respective partnerships rather than catching up on frontier capability themselves.

References

1. Anthropic. "Introducing Claude Opus 4.7." anthropic.com/news/claude-opus-4-7, April 16, 2026.
2. Anthropic. "Responsible Scaling Policy Updates." anthropic.com/rsp-updates.
3. Anthropic. "Anthropic acquires Stainless." anthropic.com/news/anthropic-acquires-stainless, May 18, 2026.
4. OpenAI. "Introducing GPT-5.5." openai.com/index/introducing-gpt-5-5, April 24, 2026.
5. OpenAI. "Introducing gpt-oss." openai.com/index/introducing-gpt-oss.
6. OpenAI. "Announcing the Stargate Project." openai.com/index/announcing-the-stargate-project, January 2025.
7. OpenAI. "OpenAI and Broadcom announce strategic collaboration." openai.com/index/openai-and-broadcom-announce-strategic-collaboration, October 13, 2025.
8. OpenAI. "Our updated Preparedness Framework." openai.com/index/updating-our-preparedness-framework, April 15, 2025.
9. Google DeepMind. "Gemini 3.1 Pro." deepmind.google/models/gemini/pro.
10. Google. "Gemma 4." blog.google/innovation-and-ai/technology/developers-tools/gemma-4, April 2026.
11. Google. "Ironwood: The first Google TPU for the age of inference." blog.google/products/google-cloud/ironwood-tpu-age-of-inference.
12. Amazon. "Amazon and Anthropic strategic collaboration." press.aboutamazon.com/2023/9/amazon-and-anthropic-announce-strategic-collaboration-to-advance-generative-ai, September 2023; expanded aboutamazon.com/news/aws/amazon-invests-additional-4-billion-anthropic-ai, November 2024.
13. Google DeepMind. "Strengthening our Frontier Safety Framework." deepmind.google/blog/strengthening-our-frontier-safety-framework, April 17, 2026.
14. AWS. "Nova 2 foundation models in Amazon Bedrock." aws.amazon.com/about-aws/whats-new/2025/12/nova-2-foundation-models-amazon-bedrock, December 2025.
15. Amazon. "Trainium3 UltraServers now available." press.aboutamazon.com/2025/12/trainium3-ultraservers-now-available, December 2025.
16. Amazon Science. "Amazon opens new AI lab in San Francisco focused on long-term research bets." amazon.science/blog/amazon-opens-new-ai-lab-in-san-francisco-focused-on-long-term-research-bets, December 9, 2024.
17. Microsoft. "Mustafa Suleyman joins Microsoft to lead Copilot." blogs.microsoft.com/blog/2024/03/19/mustafa-suleyman-deepmind-and-inflection-co-founder-joins-microsoft-to-lead-copilot, March 19, 2024.
18. Microsoft. "Introducing MAI-Transcribe-1, MAI-Voice-1, and MAI-Image-2 in Microsoft Foundry." techcommunity.microsoft.com/blog/azure-ai-foundry-blog/introducing-mai-transcribe-1-mai-voice-1-and-mai-image-2-in-microsoft-foundry/4507787.
19. Microsoft. "Maia 200: the AI accelerator built for inference." blogs.microsoft.com/blog/2026/01/26/maia-200-the-ai-accelerator-built-for-inference, January 26, 2026.

20. Microsoft. "An AI companion for everyone." blogs.microsoft.com/blog/2024/10/01/an-ai-companion-for-everyone, October 2024.
21. Meta AI. "The Llama 4 herd: The beginning of a new era of natively multimodal AI innovation." ai.meta.com/blog/llama-4-multimodal-intelligence, April 5, 2025.
22. Meta AI. "Introducing Muse Spark: Scaling Towards Personal Superintelligence." ai.meta.com/blog/introducing-muse-spark-msl, April 8, 2026.
23. Meta. "Expanding Meta's custom silicon to power our AI workloads." about.fb.com/news/2026/03/expanding-metas-custom-silicon-to-power-our-ai-workloads, March 2026.
24. Meta. "Meta Announces Joint Venture With Funds Managed by Blue Owl Capital to Develop Hyperion Data Center." about.fb.com/news/2025/10/meta-blue-owl-capital-develop-hyperion-data-center, October 21, 2025.
25. Meta. "Personal superintelligence for everyone." about.fb.com/news/2025/07/personal-superintelligence-for-everyone, July 30, 2025.
26. OpenAI. "Deliberative Alignment." openai.com/index/deliberative-alignment.
27. Amazon. "Amazon and Anthropic deepen strategic collaboration." aboutamazon.com/news/aws/amazon-invests-additional-4-billion-anthropic-ai, November 22, 2024.
28. Microsoft. "The next chapter of the Microsoft-OpenAI partnership." blogs.microsoft.com/blog/2025/10/28/the-next-chapter-of-the-microsoft-openai-partnership, October 28, 2025.
29. Meta. "Alexandr Wang, Chief AI Officer." www.meta.com/about/leadership/alexandr-wang/.
30. Meta Platforms, Inc. Q4 2025 Form 8-K, Exhibit 99.1. sec.gov/Archives/edgar/data/0001326801/000162828026003832/meta-12312025xexhibit991.htm.
31. arXiv. API documentation. info.arxiv.org/help/api/index.html.
32. OpenReview. openreview.net and API docs docs.openreview.net.
33. OpenAlex. openalex.org, API at api.openalex.org.
34. J. Priem, H. Piwowar, R. Orr. "OpenAlex: A fully-open index of scholarly works, authors, venues, institutions, and concepts." arXiv:2205.01833, 2022.
35. J. Tang, J. Zhang, L. Yao, J. Li, L. Zhang, Z. Su. "ArnetMiner: Extraction and mining of academic social networks." KDD 2008. dl.acm.org/doi/10.1145/1401890.1402008.
36. V. Prabhakaran, W. L. Hamilton, D. McFarland, D. Jurafsky. "Predicting the rise and fall of scientific topics from trends in their rhetorical framing." ACL 2016. aclanthology.org/P16-1111.
37. J. Leike on X. x.com/janleike/status/1795497960509448617, May 28, 2024.
38. J. Schulman on X. x.com/johnschulman2/status/1820610863499509855, August 5, 2024.
39. I. Sutskever on X. x.com/ilyasut/status/1803472978753303014, June 19, 2024; Safe Superintelligence Inc. ssi.inc.
40. Adept. "An update from Adept." adept.ai/blog/adept-update/, June 28, 2024.
41. Amazon. "Amazon hires from AI robotics startup Covariant, licenses its robotic foundation models." aboutamazon.com/news/company-news/amazon-covariant-ai-robots, August 30, 2024.
42. Character.AI. "Our Next Phase of Growth." blog.character.ai/our-next-phase-of-growth/, August 2, 2024.
43. Scale AI. "Scale AI announces next phase of company evolution." scale.com/blog/scale-ai-announces-next-phase-of-company-evolution, June 12, 2025.